

# Try 36 easy

May 6, 2025

```
[ ]: import os
import random
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import backend as K
from tensorflow.keras import layers, Model
from tensorflow.keras.utils import plot_model
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, auc, accuracy_score
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, \
    ModelCheckpoint
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

# Seeds
random.seed(42)
np.random.seed(42)
tf.random.set_seed(42)

# Configuration
IMG_SIZE = (128, 128)
BATCH_SIZE = 32
EPOCHS = 150
DATA_PATH = "/kaggle/input/socofing/SOCOFinF/Real/"
ALTER_EASY_PATH = "/kaggle/input/socofing/SOCOFinF/Altered/Altered-Easy/"

# Data Loading and Preprocessing
def load_fingerprint_data(data_path):
    images = []
    labels = []

    for filename in os.listdir(data_path):
        if filename.endswith(".BMP"):
            parts = filename.split('__')
            person_id = parts[0]
            finger_info = parts[1].split('_')
```

```

        hand = finger_info[1]
        finger_type = finger_info[2]
        label = f"{person_id}_{hand}_{finger_type}"

        img_path = os.path.join(data_path, filename)
        img = load_img(img_path, color_mode='grayscale',
        ↪target_size=IMG_SIZE)
        img = img_to_array(img).astype('float32') / 255.0
        images.append(img)
        labels.append(label)

    images = np.array(images)
    labels = np.array(labels)

    if images.ndim != 4:
        raise ValueError(f"Expected 4D array (num_samples, height, width,
        ↪channels), got {images.shape}")

    return images, labels

# Load both datasets
images_real, labels_real = load_fingerprint_data(DATA_PATH)
images_alter_easy, labels_alter_easy = load_fingerprint_data(ALTER_EASY_PATH)
print(f"Loaded {len(images_real)} images from Real dataset")
print(f"Loaded {len(images_alter_easy)} images from Alter-Easy dataset")
print(f"Image shape: {images_real[0].shape}")

# Pair Generation
def create_pairs(images_real, labels_real, images_alter_easy,
        ↪labels_alter_easy):
    label_to_real_image = {}
    for img, label in zip(images_real, labels_real):
        label_to_real_image[label] = img

    label_to_alter_images = {}
    for img, label in zip(images_alter_easy, labels_alter_easy):
        if label not in label_to_alter_images:
            label_to_alter_images[label] = []
        label_to_alter_images[label].append(img)

    # Generate positive pairs
    positive_pairs = []
    for label in label_to_real_image:
        if label in label_to_alter_images:
            real_img = label_to_real_image[label]
            for alter_img in label_to_alter_images[label]:
                positive_pairs.append([real_img, alter_img])

```

```

# Generate negative pairs
negative_pairs = []
num_positive = len(positive_pairs)

common_labels = list(set(label_to_real_image.keys()) &
↳set(label_to_alter_images.keys()))

while len(negative_pairs) < num_positive:
    label1 = random.choice(common_labels)
    label2 = random.choice(common_labels)
    if label1 != label2:
        real_img = label_to_real_image[label1]
        alter_img = random.choice(label_to_alter_images[label2])
        negative_pairs.append([real_img, alter_img])

# Combine and shuffle pairs
pairs = positive_pairs + negative_pairs
pair_labels = [1] * len(positive_pairs) + [0] * len(negative_pairs)

indices = np.arange(len(pairs))
np.random.shuffle(indices)
pairs = np.array(pairs)[indices]
pair_labels = np.array(pair_labels)[indices]

return pairs, pair_labels

pairs, labels = create_pairs(images_real, labels_real, images_alter_easy,
↳labels_alter_easy)
print(f"Generated {len(pairs)} pairs")

# Train/Validation/Test Split
X_train, X_temp, y_train, y_temp = train_test_split(
    pairs, labels, train_size=0.7, test_size=0.3, random_state=42,
↳stratify=labels
)

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, train_size=0.5, test_size=0.5, random_state=42,
↳stratify=y_temp
)

print(f"Training pairs: {len(X_train)}")
print(f"Validation pairs: {len(X_val)}")
print(f"Test pairs: {len(X_test)}")
print(f"Training class distribution: {np.bincount(y_train)}")
print(f"Validation class distribution: {np.bincount(y_val)}")

```

```
print(f"Test class distribution: {np.bincount(y_test)}")
```

```
2025-05-05 22:28:18.180481: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
E0000 00:00:1746484098.202392      436 cuda_dnn.cc:8310] Unable to register cuDNN
factory: Attempting to register factory for plugin cuDNN when one has already
been registered
E0000 00:00:1746484098.209488      436 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered

Loaded 6000 images from Real dataset
Loaded 17931 images from Alter-Easy dataset
Image shape: (128, 128, 1)

I0000 00:00:1746484151.365380      436 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 15513 MB memory: -> device:
0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute capability: 6.0

Generated 35862 pairs
Training pairs: 25103
Validation pairs: 5379
Test pairs: 5380
Training class distribution: [12551 12552]
Validation class distribution: [2690 2689]
Test class distribution: [2690 2690]
```

```
[2]: # Siamese Network Architecture
def create_embedding_network(input_shape):
    inputs = layers.Input(shape=input_shape)
    x = layers.Conv2D(32, (5, 5), activation='relu', kernel_regularizer=tf.
↳keras.regularizers.l2(0.01))(inputs)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.BatchNormalization()(x)

    x = layers.Conv2D(64, (3, 3), activation='relu', kernel_regularizer=tf.
↳keras.regularizers.l2(0.01))(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.BatchNormalization()(x)

    x = layers.Conv2D(128, (3, 3), activation='relu', kernel_regularizer=tf.
↳keras.regularizers.l2(0.01))(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.BatchNormalization()(x)
```

```

    x = layers.Conv2D(256, (3, 3), activation='relu', kernel_regularizer=tf.
↳keras.regularizers.l2(0.01))(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.BatchNormalization()(x)

    x = layers.Conv2D(512, (3, 3), activation='relu', kernel_regularizer=tf.
↳keras.regularizers.l2(0.01))(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.BatchNormalization()(x)

    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dense(256, activation='relu', kernel_regularizer=tf.keras.
↳regularizers.l2(0.01))(x)
    x = layers.Dropout(0.35)(x)
    x = layers.Dense(256, activation=None)(x)

    return Model(inputs, x)

def build_siamese_model(input_shape):
    input_a = layers.Input(shape=input_shape)
    input_b = layers.Input(shape=input_shape)

    embedding_network = create_embedding_network(input_shape)

    embedding_a = embedding_network(input_a)
    embedding_b = embedding_network(input_b)

    distance = layers.Lambda(
        lambda embeddings: tf.abs(embeddings[0] - embeddings[1])
    )([embedding_a, embedding_b])

    output = layers.Dense(1, activation='sigmoid')(distance)

    return Model(inputs=[input_a, input_b], outputs=output)

input_shape = IMG_SIZE + (1,)
siamese_model = build_siamese_model(input_shape)
siamese_model.summary()

# Compile the model
siamese_model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy'])

# Training
early_stopping = EarlyStopping(monitor='val_loss', patience=10,
↳restore_best_weights=True)

```

```

reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=5,
    ↪min_lr=1e-6)
checkpoint = ModelCheckpoint('best_model.keras', monitor='val_loss',
    ↪save_best_only=True)

```

Model: "functional\_1"

Layer (type)	Output Shape	Param #	Connected
↪to			
input_layer (InputLayer)	(None, 128, 128, 1)	0	-
↪			
input_layer_1	(None, 128, 128, 1)	0	-
↪			
(InputLayer)			
↪			
functional (Functional)	(None, 256)	1,769,600	
↪input_layer[0][0],			
↪input_layer_1[0][0]			
lambda (Lambda)	(None, 256)	0	
↪functional[0][0],			
↪functional[1][0]			
dense_2 (Dense)	(None, 1)	257	
↪lambda[0][0]			

Total params: 1,769,857 (6.75 MB)

Trainable params: 1,767,873 (6.74 MB)

Non-trainable params: 1,984 (7.75 KB)

```

[3]: history = siamese_model.fit(
    [X_train[:, 0], X_train[:, 1]],
    y_train,
    validation_data=(X_val[:, 0], X_val[:, 1], y_val),
    batch_size=BATCH_SIZE,

```

```

    epochs=EPOCHS,
    callbacks=[early_stopping, reduce_lr, checkpoint]
)

K.clear_session()
tf.compat.v1.reset_default_graph()

# Training Curves
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.savefig('/kaggle/working/training_curves.png')
plt.show()

# Model Evaluation
def evaluate_model(model, X, y, set_name=''):
    batch_size = 8
    y_pred = []
    for i in range(0, len(X), batch_size):
        batch_X = X[i:i + batch_size]
        batch_pred = model.predict([batch_X[:, 0], batch_X[:, 1]],
        ↪batch_size=batch_size, verbose=0)
        y_pred.extend(batch_pred)
    y_pred = np.array(y_pred)
    fpr, tpr, thresholds = roc_curve(y, y_pred)
    roc_auc = auc(fpr, tpr)

    plt.figure()
    plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area =
    ↪{roc_auc:.2f})')
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')

```

```

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'Receiver Operating Characteristic ({set_name})')
plt.legend(loc="lower right")
plt.savefig(f'/kaggle/working/roc_curve_{set_name.lower()}.png')
plt.show()

return roc_auc

print("Training Evaluation:")
train_auc = evaluate_model(siamese_model, X_train, y_train, 'Training')

print("\nValidation Evaluation:")
val_auc = evaluate_model(siamese_model, X_val, y_val, 'Validation')

print("\nTest Evaluation:")
test_auc = evaluate_model(siamese_model, X_test, y_test, 'Test')

train_acc = history.history['accuracy'][-1]
val_acc = history.history['val_accuracy'][-1]
test_acc = siamese_model.evaluate([X_test[:, 0], X_test[:, 1]], y_test)[1]

print(f"\nFinal Metrics:")
print(f"Training Accuracy: {train_acc:.4f} | AUC: {train_auc:.4f}")
print(f"Validation Accuracy: {val_acc:.4f} | AUC: {val_auc:.4f}")
print(f"Test Accuracy: {test_acc:.4f} | AUC: {test_auc:.4f}")

# Real Photo Evaluation
def visualize_predictions(model, X, y, num_samples=5):
    indices = np.random.choice(len(X), num_samples)
    sample_pairs = X[indices]
    sample_labels = y[indices]

    predictions = model.predict([sample_pairs[:, 0], sample_pairs[:, 1]])

    plt.figure(figsize=(15, 5))
    for i in range(num_samples):
        plt.subplot(2, num_samples, i+1)
        plt.imshow(sample_pairs[i][0].squeeze(), cmap='gray')
        plt.title(f"Label: {sample_labels[i]}\nPred: {predictions[i][0]:.2f}")
        plt.axis('off')

        plt.subplot(2, num_samples, i+1+num_samples)
        plt.imshow(sample_pairs[i][1].squeeze(), cmap='gray')
        plt.axis('off')

    plt.tight_layout()

```



```
plt.show()

print("Sample Test Predictions:")
visualize_predictions(siamese_model, X_test, y_test)
```

Epoch 1/150

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1746484172.874423 468 service.cc:148] XLA service 0x7a4e8c01d990 initialized for platform CUDA (this does not guarantee that XLA will be used).

Devices:

I0000 00:00:1746484172.874825 468 service.cc:156] StreamExecutor device (0): Tesla P100-PCIE-16GB, Compute Capability 6.0

I0000 00:00:1746484173.625019 468 cuda\_dnn.cc:529] Loaded cuDNN version 90300

7/785 17s 22ms/step - accuracy:  
0.5496 - loss: 10.5390

I0000 00:00:1746484179.288907 468 device\_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

785/785 41s 35ms/step -  
accuracy: 0.8686 - loss: 3.8498 - val\_accuracy: 0.7747 - val\_loss: 0.7871 -  
learning\_rate: 0.0010

Epoch 2/150

785/785 18s 22ms/step -  
accuracy: 0.9237 - loss: 0.5956 - val\_accuracy: 0.8991 - val\_loss: 0.7238 -  
learning\_rate: 0.0010

Epoch 3/150

785/785 17s 22ms/step -  
accuracy: 0.9330 - loss: 0.5545 - val\_accuracy: 0.8241 - val\_loss: 0.9546 -  
learning\_rate: 0.0010

Epoch 4/150

785/785 17s 22ms/step -  
accuracy: 0.9287 - loss: 0.6490 - val\_accuracy: 0.7816 - val\_loss: 0.9169 -  
learning\_rate: 0.0010

Epoch 5/150

785/785 18s 22ms/step -  
accuracy: 0.9281 - loss: 0.6106 - val\_accuracy: 0.9214 - val\_loss: 0.6560 -  
learning\_rate: 0.0010

Epoch 6/150

785/785 17s 22ms/step -  
accuracy: 0.9429 - loss: 0.4909 - val\_accuracy: 0.9201 - val\_loss: 0.5510 -  
learning\_rate: 0.0010

Epoch 7/150

785/785 17s 22ms/step -  
accuracy: 0.9481 - loss: 0.4764 - val\_accuracy: 0.9124 - val\_loss: 0.6225 -  
learning\_rate: 0.0010

Epoch 8/150  
785/785 17s 22ms/step -  
accuracy: 0.9415 - loss: 0.5362 - val\_accuracy: 0.9020 - val\_loss: 0.6076 -  
learning\_rate: 0.0010  
Epoch 9/150  
785/785 17s 22ms/step -  
accuracy: 0.9441 - loss: 0.4999 - val\_accuracy: 0.9334 - val\_loss: 0.4930 -  
learning\_rate: 0.0010  
Epoch 10/150  
785/785 17s 22ms/step -  
accuracy: 0.9479 - loss: 0.4553 - val\_accuracy: 0.9167 - val\_loss: 0.4833 -  
learning\_rate: 0.0010  
Epoch 11/150  
785/785 17s 22ms/step -  
accuracy: 0.9488 - loss: 0.4634 - val\_accuracy: 0.9855 - val\_loss: 0.4319 -  
learning\_rate: 0.0010  
Epoch 12/150  
785/785 17s 22ms/step -  
accuracy: 0.9509 - loss: 0.4602 - val\_accuracy: 0.9413 - val\_loss: 0.4734 -  
learning\_rate: 0.0010  
Epoch 13/150  
785/785 17s 22ms/step -  
accuracy: 0.9505 - loss: 0.4511 - val\_accuracy: 0.9446 - val\_loss: 0.4268 -  
learning\_rate: 0.0010  
Epoch 14/150  
785/785 17s 22ms/step -  
accuracy: 0.9515 - loss: 0.4256 - val\_accuracy: 0.9836 - val\_loss: 0.3512 -  
learning\_rate: 0.0010  
Epoch 15/150  
785/785 17s 22ms/step -  
accuracy: 0.9492 - loss: 0.4365 - val\_accuracy: 0.9807 - val\_loss: 0.3460 -  
learning\_rate: 0.0010  
Epoch 16/150  
785/785 17s 22ms/step -  
accuracy: 0.9567 - loss: 0.3981 - val\_accuracy: 0.9823 - val\_loss: 0.3813 -  
learning\_rate: 0.0010  
Epoch 17/150  
785/785 17s 22ms/step -  
accuracy: 0.9582 - loss: 0.3835 - val\_accuracy: 0.9221 - val\_loss: 0.6166 -  
learning\_rate: 0.0010  
Epoch 18/150  
785/785 17s 22ms/step -  
accuracy: 0.9568 - loss: 0.3888 - val\_accuracy: 0.9842 - val\_loss: 0.3666 -  
learning\_rate: 0.0010  
Epoch 19/150  
785/785 17s 22ms/step -  
accuracy: 0.9560 - loss: 0.3765 - val\_accuracy: 0.9822 - val\_loss: 0.3903 -  
learning\_rate: 0.0010

Epoch 20/150  
785/785 17s 22ms/step -  
accuracy: 0.9547 - loss: 0.3963 - val\_accuracy: 0.9175 - val\_loss: 0.5846 -  
learning\_rate: 0.0010

Epoch 21/150  
785/785 17s 22ms/step -  
accuracy: 0.9635 - loss: 0.3134 - val\_accuracy: 0.9734 - val\_loss: 0.1937 -  
learning\_rate: 2.0000e-04

Epoch 22/150  
785/785 17s 22ms/step -  
accuracy: 0.9724 - loss: 0.1639 - val\_accuracy: 0.9630 - val\_loss: 0.1995 -  
learning\_rate: 2.0000e-04

Epoch 23/150  
785/785 17s 22ms/step -  
accuracy: 0.9696 - loss: 0.1594 - val\_accuracy: 0.9740 - val\_loss: 0.1784 -  
learning\_rate: 2.0000e-04

Epoch 24/150  
785/785 17s 22ms/step -  
accuracy: 0.9715 - loss: 0.1494 - val\_accuracy: 0.9820 - val\_loss: 0.1538 -  
learning\_rate: 2.0000e-04

Epoch 25/150  
785/785 17s 22ms/step -  
accuracy: 0.9742 - loss: 0.1355 - val\_accuracy: 0.9693 - val\_loss: 0.1796 -  
learning\_rate: 2.0000e-04

Epoch 26/150  
785/785 17s 22ms/step -  
accuracy: 0.9745 - loss: 0.1383 - val\_accuracy: 0.9686 - val\_loss: 0.1730 -  
learning\_rate: 2.0000e-04

Epoch 27/150  
785/785 17s 22ms/step -  
accuracy: 0.9727 - loss: 0.1414 - val\_accuracy: 0.9894 - val\_loss: 0.1288 -  
learning\_rate: 2.0000e-04

Epoch 28/150  
785/785 17s 22ms/step -  
accuracy: 0.9783 - loss: 0.1283 - val\_accuracy: 0.9872 - val\_loss: 0.1395 -  
learning\_rate: 2.0000e-04

Epoch 29/150  
785/785 17s 22ms/step -  
accuracy: 0.9761 - loss: 0.1311 - val\_accuracy: 0.9894 - val\_loss: 0.1448 -  
learning\_rate: 2.0000e-04

Epoch 30/150  
785/785 17s 22ms/step -  
accuracy: 0.9767 - loss: 0.1376 - val\_accuracy: 0.9797 - val\_loss: 0.1499 -  
learning\_rate: 2.0000e-04

Epoch 31/150  
785/785 17s 22ms/step -  
accuracy: 0.9754 - loss: 0.1296 - val\_accuracy: 0.9708 - val\_loss: 0.1703 -  
learning\_rate: 2.0000e-04

Epoch 32/150  
785/785 17s 22ms/step -  
accuracy: 0.9678 - loss: 0.1679 - val\_accuracy: 0.9892 - val\_loss: 0.1396 -  
learning\_rate: 2.0000e-04

Epoch 33/150  
785/785 17s 22ms/step -  
accuracy: 0.9792 - loss: 0.1235 - val\_accuracy: 0.9870 - val\_loss: 0.1228 -  
learning\_rate: 4.0000e-05

Epoch 34/150  
785/785 17s 22ms/step -  
accuracy: 0.9832 - loss: 0.0960 - val\_accuracy: 0.9883 - val\_loss: 0.1103 -  
learning\_rate: 4.0000e-05

Epoch 35/150  
785/785 17s 22ms/step -  
accuracy: 0.9849 - loss: 0.0858 - val\_accuracy: 0.9894 - val\_loss: 0.1020 -  
learning\_rate: 4.0000e-05

Epoch 36/150  
785/785 17s 22ms/step -  
accuracy: 0.9832 - loss: 0.0815 - val\_accuracy: 0.9872 - val\_loss: 0.1055 -  
learning\_rate: 4.0000e-05

Epoch 37/150  
785/785 17s 22ms/step -  
accuracy: 0.9845 - loss: 0.0765 - val\_accuracy: 0.9861 - val\_loss: 0.1095 -  
learning\_rate: 4.0000e-05

Epoch 38/150  
785/785 17s 22ms/step -  
accuracy: 0.9849 - loss: 0.0766 - val\_accuracy: 0.9822 - val\_loss: 0.1141 -  
learning\_rate: 4.0000e-05

Epoch 39/150  
785/785 17s 22ms/step -  
accuracy: 0.9857 - loss: 0.0735 - val\_accuracy: 0.9874 - val\_loss: 0.1038 -  
learning\_rate: 4.0000e-05

Epoch 40/150  
785/785 17s 22ms/step -  
accuracy: 0.9876 - loss: 0.0677 - val\_accuracy: 0.9853 - val\_loss: 0.1079 -  
learning\_rate: 4.0000e-05

Epoch 41/150  
785/785 17s 22ms/step -  
accuracy: 0.9896 - loss: 0.0632 - val\_accuracy: 0.9849 - val\_loss: 0.1081 -  
learning\_rate: 8.0000e-06

Epoch 42/150  
785/785 17s 22ms/step -  
accuracy: 0.9912 - loss: 0.0584 - val\_accuracy: 0.9853 - val\_loss: 0.1040 -  
learning\_rate: 8.0000e-06

Epoch 43/150  
785/785 17s 22ms/step -  
accuracy: 0.9909 - loss: 0.0574 - val\_accuracy: 0.9846 - val\_loss: 0.1064 -  
learning\_rate: 8.0000e-06

Epoch 44/150

785/785 17s 22ms/step -

accuracy: 0.9915 - loss: 0.0563 - val\_accuracy: 0.9848 - val\_loss: 0.1059 -


learning\_rate: 8.0000e-06

Epoch 45/150


785/785 17s 22ms/step -

accuracy: 0.9918 - loss: 0.0539 - val\_accuracy: 0.9853 - val\_loss: 0.1068 -


learning\_rate: 8.0000e-06




Training Evaluation:



Validation Evaluation:



Test Evaluation:



169/169                      2s 10ms/step -  
accuracy: 0.9928 - loss: 0.0991

Final Metrics:

Training Accuracy: 0.9922 | AUC: 1.0000

Validation Accuracy: 0.9853 | AUC: 0.9999

Test Accuracy: 0.9913 | AUC: 0.9999

Sample Test Predictions:

1/1                      1s 601ms/step



